

FAT (File Allocation Table) File System Tutorial

<http://students.cs.byu.edu/~cs345ta/lectures/seamons/FAT.pdf>

1

FAT File System

- Originated in the late 1970s and early 1980s
- The file system supported by MS-DOS
- Originally developed as a simple file system suitable for floppy disk drives less than 500K in size
- Currently there are three FAT file system types: FAT12, FAT16, FAT32
 - FATnn – nn is the number of bits in each entry in the FAT structure on the disk



2

FAT File System Volume

- Four basic regions
 - 0 – Reserved Region
 - First disk sector
 - Boot sector, reserved sector, 0th sector
 - 1 – FAT Region
 - Always 2 FATs for redundancy
 - 2 – Root Directory Region (non FAT32)
 - Fixed size
 - 3 – File and Directory Data Region
 - Files are stored in clusters (cluster = N sectors)
 - Directories are specially formatted files





















































3

Boot Sector (FAT12,FAT16)

```
BYTE jumpBoot[3]           // jump instruction to boot code
BYTE SysName[8]            // name string (MSDOS5.0)
WORD BytesPerSector        // 512, 1024, 2048, 4096
BYTE SectorsPerCluster     // 1, 2, 4, 8, 16, 32, 64
WORD ReservedSectors       // 1 (for FAT12 and FAT16)
BYTE FATcount              // # of FAT data structures (always 2)
WORD MaxRootEntries        // # of 32-byte directory entries in root dir
WORD TotalSectors1         // Total sectors in the volume
BYTE MediaDescriptor       // Removable media is 0xF0
WORD SectorsPerFAT         // Sectors in ONE FAT
WORD SectorsPerTrack       //
WORD HeadCount
DWORD HiddenSectors
DWORD TotalSectors2
BYTE DriveNumber
BYTE Reserved1
BYTE ExtBootSignature
DWORD VolumeSerial         // Volume serial number
BYTE VolumeLabel[11]      // Volume label
BYTE Reserved2[8]
```

4

Disk Structure

| | | |
|---|----------------|--|
|  | Boot Sector | 0 |
|  | FAT 1 | 1 thru SectorsPerFAT |
|  | FAT 2 | (SectorsPerFAT+1) thru (2*SectorsPerFAT) |
|  | Root Directory | (2*SectorsPerFAT + 1) thru (2*SectorsPerFAT + RootDirSectors) |
|  | | |
|  | File clusters | (2*SectorsPerFAT + 1 + RootDirSectors) thru TotalSectors - 1 |
|  | | |
|  | | |
|  | | |
|  | | |
|  | | |
|  | | |
|  | | |
|  | | |
|  | | |
|  | | |
|  | | |
|  | | |
|  | | |
|  | | |
|  | | |
|  | | |
|  | | |
|  | | |
|  | | |
|  | | |
|  | | |
|  | | |
|  | | |
|  | | |
|  | | |
|  | | |
|  | | |
|  | | |
|  | | |
|  | | |
|  | | |
|  | | |
|  | | |
|  | | |
|  | | |
|  | | |
|  | | |
|  | | |
|  | | |
|  | | |
|  | | |
|  | | |
|  | | |
|  | | |
|  | | |
|  | | |


5

FAT Region

- There are FATcount (2) FATs in the volume
 - Keep both tables consistent
- Each FAT is SectorsPerFat sectors long
 - FAT12 lab – 9
- Each FAT is a singly linked list of the clusters of a file
- All linking is done using logical cluster numbers
- The first data cluster number is 2
- The first sector of cluster 2
 - $\text{RootDirSectors} = ((\text{MaxRootEntries} * 32) + (\text{BytesPerSector} - 1)) / \text{BytesPerSector}$
 - $\text{FirstDataSector} = \text{ReservedSectors} + (\text{FATcount} * \text{SectorsPerFAT}) + \text{RootDirSectors}$
- The first sector of any cluster N
 - $\text{FirstSectorOfCluster} = ((N - 2) * \text{SectorsPerCluster}) + \text{FirstDataSector}$

6

FAT File System



7

FAT12

- Each FAT12 entry is 12 bits
 - When designed, space was tight
 - Pack 2 entries into 3 bytes
 - 4096 possible clusters
 - If a sector is 512 bytes and cluster = 1 sector, can represent 2 Mb of data
- First 2 entries in the FAT are reserved
 - Don't use
 - Logical cluster numbers begin at 2
- FAT12 entry values

| | |
|----------------|--------------------------------------|
| – 0 | Unused cluster |
| – 0xFF0-0xFF6 | Reserved cluster |
| – 0xFF7 | Bad cluster |
| – 0xFF8-0xFFFF | End Of Clusterchain mark (EOC) <EOF> |
| Other | Next cluster in file |

8


FAT12 entry packing

- Bytes $3x$ to $3x+2$
 - contain FAT entries for logical clusters $2x$ and $2x+1$
- Byte $3x$
 - 8 least sig bits of entry $2x$
- Byte $3x+1$
 - 4 least sig bits are the 4 most sig bits of entry $2x$
 - 4 most sig bits are the 4 least sig bits of entry $2x+1$
- Byte $3x+2$
 - 8 most sig bits of entry $2x+1$



9

FAT12 entry packing



- Example FAT:
 - F0 FF FF 03 40 00 05 60 00 07 80 00 09 F0 FF
- First 3 bytes are first 2 entries → reserved
- Suppose first file cluster is 2 → bytes 3 and 4
- Converting entries 2 and 3
 - 03 40 00 → 003 and 004
- FF8 – FFF indicates EOF



10

Getting a FAT12 Entry

- Determine byte offset in the FAT
 - $\text{FatIndex} = (\text{LogicalCluster} * 3) / 2$
 - $\text{FatIndex} + 1$
- Determine if an Even or Odd Entry
- Split LogicalCluster
 - Even \rightarrow 4 most sig, 8 least sig
 - Odd \rightarrow 8 most sig, 4 least sig
- Build next Logical Cluster number
 - Start with zero
 - Copy most sig bits and shift
 - Or in least sig bits

11

Get FAT12 entry (bytes)

```
int getCluster(int LogicalCluster)
{
    sector = 0
    // Convert LogicalCluster to byte location
    FatIndex = (LogicalCluster * 3) / 2
    if (LogicalCluster is even)
        //split FatIndex + 1 for 4 most sig bits
        sector = FAT[FatIndex + 1] & 0x0f
        sector = sector << 8
        sector = sector | FAT[FatIndex]
    else
        //split FatIndex for 4 least sig bits
        sector = FAT[FatIndex+1]
        sector = sector << 4
        sector = sector | ((FAT[FatIndex] & 0xf0) >> 4)
}
```

12

Setting a FAT12 Entry

- Determine byte offsets in the FAT
 - $\text{FatIndex} = (\text{LogicalCluster} * 3) / 2$
 - $\text{FatIndex} + 1$
- Determine if an Even or Odd entry
- Split LogicalCluster
 - Even \rightarrow 4 most sig, 8 least sig
 - Odd \rightarrow 8 most sig, 4 least sig
- Set FAT12 table entry

13

Set FAT12 entry (bytes)

```
void setFatEntry(int LogicalCluster, int value)
{
    // Convert LogicalCluster to byte location
    FatIndex = (LogicalCluster * 3) / 2
    if (LogicalCluster is even)
        //split FatIndex + 1 for 4 most sig bits
        FAT[FatIndex+1] &= 0xf0
        FAT[FatIndex+1] |= (value & 0xf00) >> 8
        FAT[FatIndex] = value & 0x0ff
    else
        //split FatIndex for 4 least sig bits
        FAT[FatIndex] &= 0x0f
        FAT[FatIndex] |= (value & 0x00f) << 4
        FAT[FatIndex+1] = (value & 0xff0) >> 4
}
```

14

Byte Ordering

- Intel and FAT byte ordering is Little Endian.
- FAT12 table in Hex
 - each byte given in Little Endian order
 - F0 FF FF 03 40 00 05 60 00 07 80 00 09 F0 FF
 - Ordering nibbles (half-byte) from least to most significant
 - 0F FF FF 30 04 00 50 06 00 70 08 00 90 0F FF
- The Little Endian arrangement of bytes makes it convenient to extract and set the entries using a byte offset to the WORD containing the entry in the table
 - An entry is the least significant (Even) or most significant (Odd) 12 bits of the WORD

15

Get FAT12 Entry (WORDS)

```
FATEntOffset = (ClusterNum * 3)/2
Fat12ClusterEntry = *((WORD*) &FATbyte[FATEntOffset]);
If(ClusterNum & 0x0001)
    // entry is the high order 3 bytes of the WORD
    Fat12ClusterEntry = Fat12ClusterEntry >> 4;
Else
    // entry is the low order 3 bytes of the WORD
    Fat12ClusterEntry = Fat12ClusterEntry & 0x0FFF;
```

16

Set FAT12 Entry (WORDS)

```
FatEntOffset = (ClusterNum * 3) / 2
if (ClusterNum & 0x0001)
    // entry is the high order 3 bytes of the WORD
    Fat12ClusterEntry = Fat12ClusterEntry << 4;
    *((WORD*) &FAT[FatEntOffset]) =
        *((WORD*) &FAT[FatEntOffset]) & 0x000F;
Else
    // entry is the low order 3 bytes of the WORD
    Fat12ClusterEntry = Fat12ClusterEntry & 0x0FFF;
    *((WORD*) &FAT[FatEntOffset]) =
        *((WORD*) &FAT[FatEntOffset]) & 0xF000;

*((WORD*) &FAT[FatEntOffset]) =
    ((*((WORD*) &FAT[FatEntOffset])) | Fat12ClusterEntry;
```



17

Typical FAT12 Floppy Disk

- 1 sector per cluster
- 1 sector reserved (Boot sector)
- 18 sectors for 2 FAT tables
- 14 sectors for root directory
- Convert logical cluster number to sector number
 - Subtract 2 from logical cluster number
 - Multiply by number of sectors per cluster (1)
 - Add result to first data area sector number (33)



18

Directories

- Array of directory entries
- Root directory has a fixed number of entries
 - 14 sectors reserved → How many entries?
 - Contiguous sectors
- Subdirectories are simply files
 - Same array of directory entries structure within each cluster (no size restriction)
- Directory entry
 - 32 bytes
 - Filename's first character is usage indicator
 - 0x00 Never been used
 - 0xe5 Used before but entry has been released

19

Directory Entry

```
BYTE Filename[8]      // Pad with spaces
BYTE Extension[3]     //
BYTE Attributes
BYTE Reserved[10]
WORD Time              // Hour*2048+Min*32+sec/2
WORD Date              // (Yr-1980)*512 + Mon*32 + day
WORD StartCluster
DWORD FileSize
```

20

File Attributes

| <u>Bit</u> | <u>Mask</u> | <u>Attribute</u> |
|------------|-------------|------------------|
| 0 | 0x01 | Read Only |
| 1 | 0x02 | Hidden |
| 2 | 0x04 | System |
| 3 | 0x08 | Volume Label |
| 4 | 0x10 | Subdirectory |
| 5 | 0x20 | Archive |
| 6 | 0x40 | Unused |
| 7 | 0x80 | Unused |

21

Suggested Lab Progression

- Read boot sector and print out
- Read root directory and print out
- Read FAT and interpret for files in root dir
 - Print out cluster list for each file
- Read in and print out a file
- Read in and print out a subdirectory
- Finish up
- Make DLL

22